

FinalReport-G39 (1)

May 29, 2024

#

Group 39 STAT 301 Final Report

####

Group Members: Alice Le, Sohbat Sandhu, Tanvi Singh, Zimo Zhang

0.1 *Predictive Analysis on Prices of Rental Apartments in USA using Major Rental Price Influencers*

0.2 Introduction

0.2.1 Data Information

We are using a [Dataset of classified for apartments for rent in United States of America](#). The dataset contains 10,000 instances of classified apartments with 22 features (listed in the table below). The data has been originally cleaned in such a way that **column** *price* and *square_feet* variables are never empty but the dataset is saved as it was created.

The data has been collected from 12 unique online sources for rental listings and has been collected from September 2019 to December 2019.

Variables	Data type	Description
id	double	Every Apartment on the classified is given a Unique Identifier
category	character	Category of the apartment classified informs us about the type of rental property
title	character	Title of the rental property listed in the Classified
body	character	Description of the rental apartment listed in the Classified
amenities	character	Amenities included with the rental apartment

Variables	Data type	Description
bathrooms	<i>character</i>	Number of bathrooms in the rental apartment
bedrooms	<i>character</i>	Number of bedrooms in the rental apartment
currency	<i>character</i>	Currency used for the price listing of the rental apartment in the Classified
fee	<i>character</i>	Additional Apartment Fee for the renting the rental apartment in the Classified
has_photo	<i>character</i>	Does the rental apartment listed in the Classified comes with a photo?
pets_allowed	<i>character</i>	Types of pets allowed in the rental apartment
price	<i>double</i>	Rental Price of Apartment
price_display	<i>character</i>	Price converted into display for reader in the classified
price_type	<i>character</i>	Price in USD of the rental apartment
square_feet	<i>double</i>	Size of rental Apartment in Square Feet as listed in the classified
address	<i>character</i>	Street Address of the rental apartment
cityname	<i>character</i>	City Location of the rental apartment
state	<i>character</i>	State/District Location of the rental apartment
latitude	<i>character</i>	Latitude Coordinates of the rental apartment
longitude	<i>character</i>	Longitude Coordinates of the rental apartment
source	<i>character</i>	Source from which the classified was taken from or the source at which the apartment was listed for rental

Variables	Data type	Description
time	<i>double</i>	Time at which the the apartment was listed for rental (in Epoch Unix Timestamp)

Table 1. Information of features in the Classified Dataset Through analysis of our dataset, we are looking to develop a predictive model that predicts rental prices in the USA.

According to Nishani (2016), the attributes that affect rental prices can be categorised into three major aspects, namely: physical attributes, locational attributes, and the amenities provided.

Based on our chosen dataset, we can split up relevant explanatory variables for our chosen response variable **price** such: * Physical attributes — **bathrooms**, **bedrooms**, **square_feet** * Locational attributes — **state** * Amenities provided — **pets_allowed**, **amenities**

Other variables are omitted due to their redundancy, due to the fact that they do not fall under the three identified categories, or they cannot be quantified in a way that can be used in this research.

The representative dataset includes information about features owners and potential tenants have at their disposal when judging the rental price of apartments from the various states and regions in the USA.

We have chosen these variables due to their importance in choosing a place to rent, along with the fact that they have varying levels that can also impact rental prices. The variation will be a good way to produce a predictor tool that accounts for many different types of observations.

Creating a predictor tool like this is crucial, especially in the midst of a housing crisis. There are a range of difficulties when trying to find a house to rent, when there are many criteria to adhere to.

For example, pet-owners would have a hard time finding a more affordable place to live, considering that it is considered a “luxury”, accompanied with a pet fee or a pet deposit. On the other hand, the location of the apartment, based on the US state, is a huge factor in indicating the rental price. Apartments located in metropolitan areas tend to be more expensive than those in areas with more vacancy Collinson (2009).

We aim to create a predictor tool that has the potential to compensate for and/or incorporate any new explanatory variables that may be added in the future. On the same note, we will be comprehensive enough to create a tool that can be used elsewhere, in order to predict rental prices elsewhere.

0.3 Methods and Results

0.3.1 Data Preparation (Import, Clean and Wrangle Data)

```
[1]: # Load useful libraries
library(tidyverse)
library(tidymodels)
library(GGally)
library(infer)
library(AER)
```

```
library(leaps)
library(dplyr)
library(reshape2)
```

```
Attaching core tidyverse packages          tidyverse
2.0.0
dplyr      1.1.3      readr      2.1.4
forcats    1.0.0      stringr    1.5.0
ggplot2     3.4.3      tibble     3.2.1
lubridate  1.9.3      tidyr      1.3.0
purrr       1.0.2
```

Conflicts

```
tidyverse_conflicts()
dplyr::filter() masks stats::filter()
dplyr::lag()    masks stats::lag()
Use the conflicted package
(<http://conflicted.r-lib.org/>) to force all conflicts to
become errors
```

```
Attaching packages          tidymodels
1.1.1
```

```
broom      1.0.5      rsample
1.2.0
dials      1.2.0      tune
1.1.2
infer      1.0.5      workflows
1.1.3
modeldata  1.2.0      workflowsets
1.0.1
parsnip    1.1.1      yardstick
1.2.0
recipes    1.0.8
```

Conflicts

```
tidymodels_conflicts()
scales::discard() masks
purrr::discard()
dplyr::filter()   masks
stats::filter()
recipes::fixed()  masks
stringr::fixed()
dplyr::lag()      masks stats::lag()
yardstick::spec() masks readr::spec()
recipes::step()   masks stats::step()
• Dig deeper into tidy modeling with R at
https://www.tmw.r.org
```

```
Registered S3 method overwritten by 'GGally':  
  method from  
    +.gg      ggplot2
```

```
Loading required package: car
```

```
Loading required package: carData
```

```
Attaching package: 'car'
```

```
The following object is masked from 'package:dplyr':
```

```
  recode
```

```
The following object is masked from 'package:purrr':
```

```
  some
```

```
Loading required package: lmtest
```

```
Loading required package: zoo
```

```
Attaching package: 'zoo'
```

```
The following objects are masked from 'package:base':
```

```
  as.Date, as.Date.numeric
```

```
Loading required package: sandwich
```

```
Loading required package: survival
```

```
Attaching package: 'reshape2'
```

```
The following object is masked from 'package:tidyr':
```

```
  smiths
```

```
[2]: # Assign dataset url which was uploaded to github from the original website for
      ↪ easier access
classified_url <- "https://raw.githubusercontent.com/vyle2003/STAT_301_Project/
      ↪ main/apartments_for_rent_classified_10K.csv"

# Read csv file and assign
apartments <- read_csv2(classified_url)
head(apartments, 1)
```

Using `"', '"` as decimal and `"'. '"` as grouping mark. Use ``read_delim()`` for more control.

Rows: 10000 Columns: 22
Column specification

Delimiter: `","`

`chr` (18): category, title, body, amenities, bathrooms, bedrooms, currency, f...
`dbl` (4): id, price, square_feet, time

Use ``spec()`` to retrieve the full column specification for this data.

Specify the column types or set ``show_col_types = FALSE`` to quiet this message.

	id	category	title
A tibble: 1 × 22	<dbl>	<chr>	<chr>
	5668626895	housing/rent/apartment	Studio apartment 2nd St NE, Uhland Terrace NE, Was

```
[3]: # Excluding variables that are unnecessary for analysis
      # excluding "id", "title", "body", "source" : variables not relevant to
      ↪ question of interest
      # excluding "category", "currency", "fee" : show negligible variation in
      ↪ values
      # excluding "has_photo": price calculated by the owner, consumers opinion
      ↪ about the photo leads to no variation in price
      # excluding "address", "cityname", "latitude", "longitude" : using location
      ↪ in a broader scope (using "state" variable)
# assigning dataframe to a new variable
apartmentsTidy <- apartments %>%
  select(-c("id" : "body", "currency", "fee", "has_photo", "address",
  ↪ "cityname", "latitude", "longitude", "source", "time"))
head(apartmentsTidy)
```

	amenities <chr>	bathrooms <chr>	bedrooms <chr>	pets_allowed <chr>	price <dbl>
A tibble: 6 × 9	null	null	0	None	790
	null	null	1	None	425
	null	1	0	None	1390
	null	1	0	None	925
	null	null	0	None	880
	Dishwasher,Elevator,Patio/Deck,Pool,Storage	1	0	null	2475

Table 2. Feature values for the relevant variables of Classified Dataset

```
[4]: # Tidy data for easier analysis

# Assigning appropriate values for the number of bathrooms (1) and bedrooms (0)
# for studio apartments
# (represented as "null" values in dataset as studio apartments)
apartmentsTidy$bathrooms <- ifelse(apartmentsTidy$bathrooms == "null", "1",
# apartmentsTidy$bathrooms)
apartmentsTidy$bedrooms <- ifelse(apartmentsTidy$bedrooms == "null", "0",
# apartmentsTidy$bedrooms)

# Transforming price of an apartment where the rent is not "Monthly" (price is
# calculated "Weekly")
apartmentsTidy$price <- ifelse((apartmentsTidy$price_type != "Monthly"),
4*apartmentsTidy$price, apartmentsTidy$price) #
# as there are 4 weeks in a month

# Calculate the number of amenities and place the number in a separate column
# named 'no_of_amenities'
apartmentsTidy <- apartmentsTidy %>%
  mutate(no_of_amenities = ifelse(amenities == "null",
0, str_count(amenities, ",") + 1)) %>%
  select(-amenities)

# Convert relevant variables with chr data types to dbl
apartmentsTidy <- apartmentsTidy %>%
  mutate(bathrooms = as.numeric(bathrooms)) %>%
  mutate(bedrooms = as.numeric(bedrooms)) %>%
  select(-c("price_display", "price_type"))

# Assign appropriate values for "null" in pets_allowed column
# (we are assuming that as there is no information about the pets policy there
# is no restriction on pets)
apartmentsTidy$pets_allowed <- ifelse(apartmentsTidy$pets_allowed == "null",
#"Yes", apartmentsTidy$pets_allowed)
apartmentsTidy$pets_allowed <- ifelse(apartmentsTidy$pets_allowed ==
#"Cats,Dogs", "Cats&Dogs", apartmentsTidy$pets_allowed)
```

```

apartmentsTidy$pets_allowed <- ifelse(apartmentsTidy$pets_allowed == "Cats",
  ↪ "OnlyCats", apartmentsTidy$pets_allowed)
apartmentsTidy$pets_allowed <- ifelse(apartmentsTidy$pets_allowed == "Dogs",
  ↪ "OnlyDogs", apartmentsTidy$pets_allowed)

# Remove outliers for better prediction performance
lower <- quantile(apartmentsTidy$price, 0.5) - 1.5*IQR(apartmentsTidy$price)
upper <- quantile(apartmentsTidy$price, 0.5) + 1.5*IQR(apartmentsTidy$price)

apartmentsTidy <- apartmentsTidy %>%
  filter((price <= upper) & (price >= lower))
head(apartmentsTidy)

```

	bathrooms	bedrooms	pets_allowed	price	square_feet	state	no_of_amenities
	<dbl>	<dbl>	<chr>	<dbl>	<dbl>	<chr>	<dbl>
A tibble: 6 × 7	1	0	None	790	101	DC	0
	1	1	None	425	106	IN	0
	1	0	None	1390	107	VA	0
	1	0	None	925	116	WA	0
	1	0	None	880	125	VA	0
	1	0	None	1800	132	CA	0

Table 3. Tidied Table

0.3.2 Exploratory Data Analysis and Visualization

```

[5]: # Plot scatterplot: price vs size of an apartment
apartmentsTidy %>% ggplot(aes(x = square_feet, y = price)) +
  geom_point() +
  labs(x = "Apartment Size (in sqft)", y = "Apartment Rental Price (in USD)")
  ↪+
  ggtitle("Scatterplot of Apartment Rental Price and Apartment Price") +
  theme(text = element_text(size = 15)) +
  scale_y_continuous(limits = c(0, 2500)) +
  scale_x_continuous(limits = c(0, 2000))

```

Warning message:

"Removed 198 rows containing missing values (`geom_point()`)."




Figure 1. Scatterplot between Apartment Size and Apartment Price

Observations: From Figure.1 we observe a possible positive correlation between the Apartment Size and Apartment Price indicating that `square_feet` is a potential input variable for Predictive Model

```
[6]: # Explore a potential problem of multicollinearity using pair plots and finding
      ↪ correlation between variables
cor_matrix <- cor(subset(apartmentsTidy, select = -c(pets_allowed, state)))
melted <- melt(cor_matrix)

corr_plot <- ggplot(melted) +
  geom_tile(aes(Var1, Var2, fill=value), colour = "black") +
```

```

    geom_text(aes(Var1, Var2, label = round(value,2)), color = "black", size = 3.5) +
    theme(axis.text.x = element_text(angle = 45, vjust = 1, hjust=1), text = element_text(size = 12, face = "bold")) +
    scale_fill_gradient2(low = "#6D9EC1", high = "#4e36b5")+
    guides(fill = guide_colourbar(barwidth = 0.5, barheight = 20)) +
    labs(x = "Variable 1", y = "Variable 2") +
    ggtitle("Correlation Matrix for the Continous Variables")
corr_plot

```

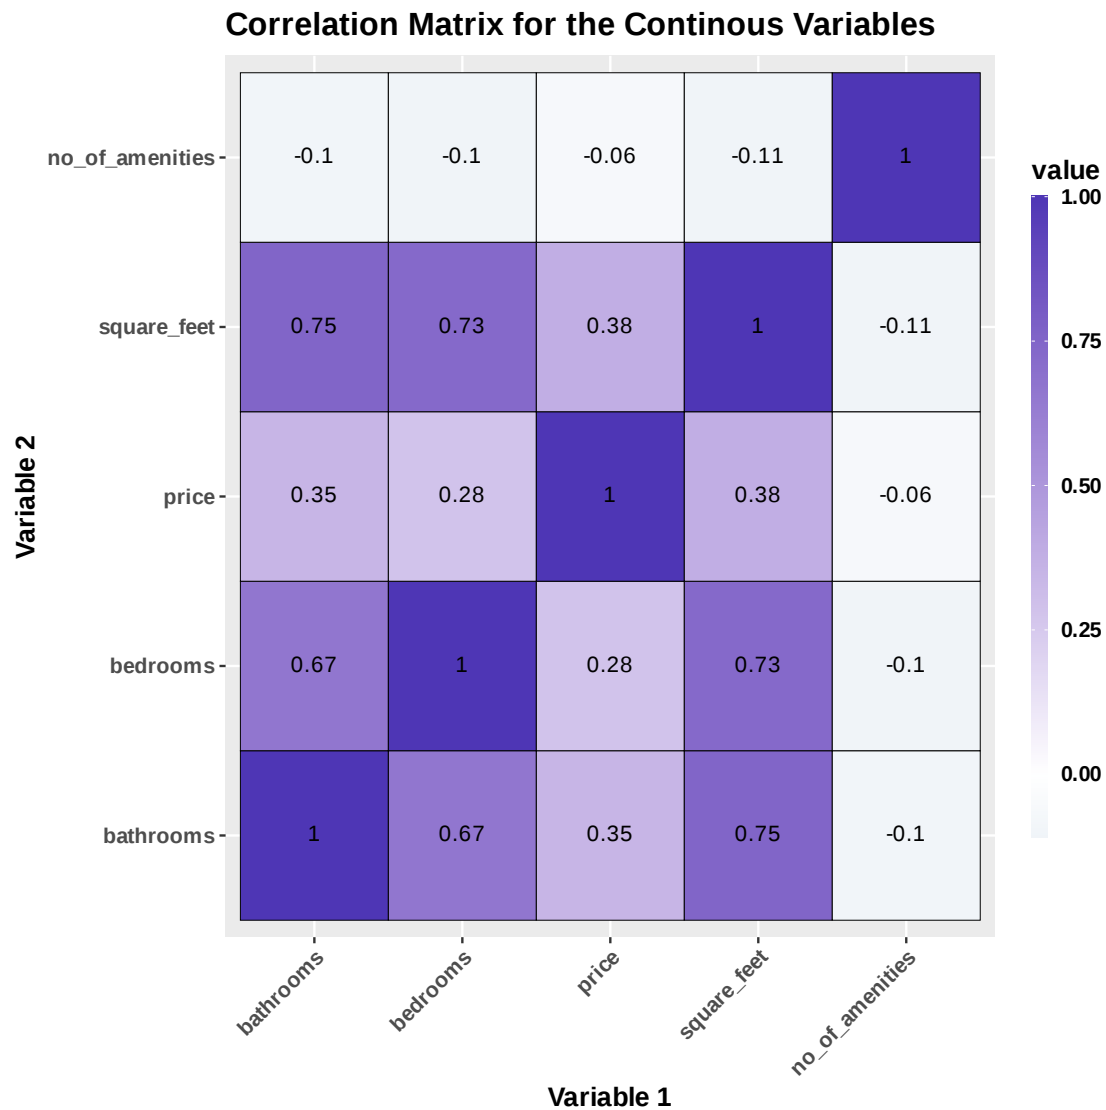


Figure 2. Heatmap for calculating correlation

Observations: From Figure.2 we observe multicollinearity between possible input vari-

ables. There is high correlation between `bathrooms` and `square_feet`, `bathrooms` and `bedrooms`, and `bedrooms` and `square_feet`. The figure also shows poor correlation between predictor variable `price` and possible input variable `no_of_amenities`, indicating `no_of_amenities` a poor choice as input variable.

```
[7]: options(repr.plot.width = 10, repr.plot.height = 5)

# Calculate the median price for each state
medians <- aggregate(price ~ state, data = apartmentsTidy, FUN = median)

# Sort the data frame by median values in increasing order
medians <- medians[order(-medians$price), ]

# Plot the median Apartment Price (scaled by 0.0001) according to each state in_
↳Descending order
apartmentsTidy %>% select(c("state", "price", "square_feet")) %>%
  group_by(state) %>% summarize(price = 0.001*median(price)) %>%
  ggplot() +
  geom_line(aes(x = reorder(state, -price), y = price, group = 1)) +
  geom_point(aes(x = reorder(state, -price), y = price, color = state)) +
  labs(x = "US State", y = "Median Apartment Rental Price for State (in_
↳USD)", color = "US State") +
  ggtitle("Line Chart of Median Apartment Rental Price for each US State") +
  theme(text = element_text(size = 10))
```




Figure 3. Line Chart of Median Apartment Rental Price for each US State

Observations: From Figure.3 we observe that Apartment Rental Price varies for each state indicating `state` as a potential input variable for our predictive model.

```
[8]: # Boxplot to explore the relationship of the categories for pets_allowed and
      ↪price
apartmentsTidy %>%
  ggplot(aes(x = pets_allowed, y = price, fill = pets_allowed)) +
  geom_boxplot() +
  labs(x = "Pets Allowed", y = "Rental Price (in USD)", fill = "Pets_
      ↪Allowed") +
  ggtitle("Boxplot of Rental Price for respective Pets Allowed Policy") +
  theme(text = element_text(size = 12))
```

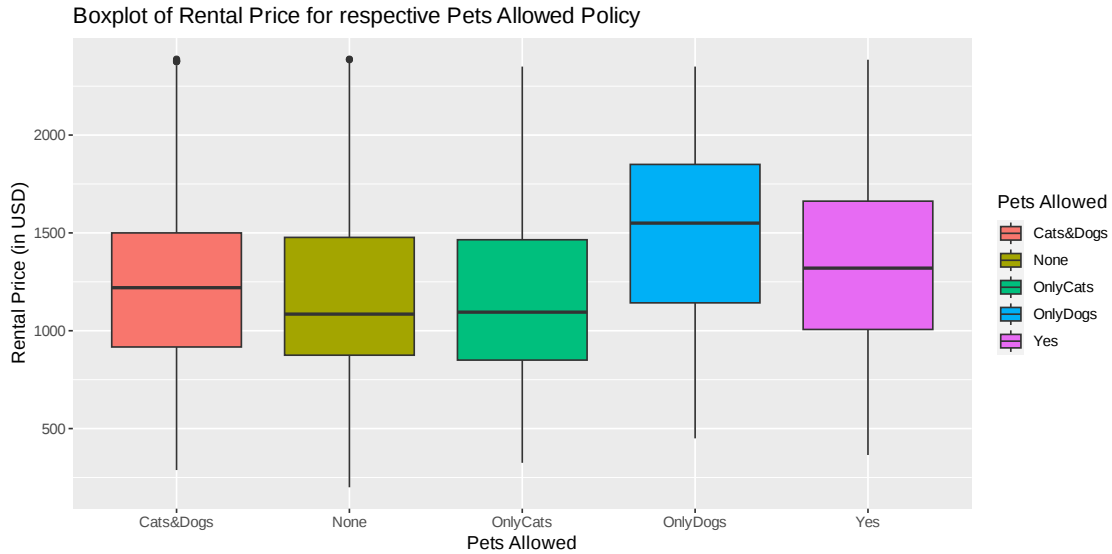


Figure 4. Boxplot of Apartment Rental Price for the respective Pet Policy

Observations: From Figure.4 we observe that Apartment Rental Price varies for each pet policy indicating `pets_allowed` as a potential input variable for our predictive model. A huge visible difference in the median apartment price can be seen for the “Only Dogs” and “Yes” category in comparison to other policies

0.3.3 Methods: Plans

As we are making a predictive model for out-of-sample predictions, we first need to split the data into two datasets, training and testing dataset. The training dataset will be used to train the regression model. To identify the most important variables or predictors that significantly contribute to a model’s predictive power, we would use **Forward Stepwise Selection** method to evaluate the linear model metrics such as Mallows’s C_p , AIC and BIC are computed with the training dataset and select the most appropriate model that minimizes the approximate **Root Mean Squared Error** from the testing dataset.

But before that we have to resolve the problems we face due to the huge number of categories in the `state` variable. To resolve this we have to compress the `state` input variable into a categorical

variable with three levels. The three levels after transformation will indicate whether the apartment lies in the state part of the West Coast, East Coast or Central Midwest Region. Similarly, the `pets_allowed` variable should also be transformed into columns for easier use of the stepwise selection algorithm.

```
[9]: # Compress state variable into regions: West Coast, East Coast and Central
      ↪Midwest
west_coast <- c("WA", "OR", "CA", "AK", "HI")
east_coast <- c("ME", "NH", "VT", "MA", "RI", "CT", "NY", "NJ", "PA", "DE",
      ↪"MD", "VA", "WV", "NC", "SC", "GA", "FL")
central_midwest <- c("OH", "IN", "MI", "IL", "WI", "MN", "IA", "MO", "KS",
      ↪"NE", "SD", "ND", "KY", "TN", "AL", "MS",
      ↪"AR", "LA", "TX", "OK", "NM", "AZ", "MT", "ID", "WY",
      ↪"CO", "UT", "NV")

# Remove "null" state categories
apartmentsTidy <- apartmentsTidy %>%
  filter(state != "null")

# function to assign region to state
assign_region <- function(state) {
  if (state %in% west_coast) {
    return("WestCoast")
  } else if (state %in% east_coast) {
    return("EastCoast")
  } else {
    return("CentralMidwest")
  }
}

# Create "region" and remove "state" variable from dataset
apartmentsTidy <- apartmentsTidy %>%
  mutate(region = sapply(state, assign_region)) %>%
  select(-state)
```

```
[10]: head(apartmentsTidy)
```

	bathrooms <dbl>	bedrooms <dbl>	pets_allowed <chr>	price <dbl>	square_feet <dbl>	no_of_amenities <dbl>	region <chr>
A tibble: 6 × 7	1	0	None	790	101	0	CentralMidw
	1	1	None	425	106	0	CentralMidw
	1	0	None	1390	107	0	EastCoast
	1	0	None	925	116	0	WestCoast
	1	0	None	880	125	0	EastCoast
	1	0	None	1800	132	0	WestCoast

Table 4. Transformed Data Stepwise selection assumes the basic linear model assumptions and this might not always give the most optimal subsets. For example, the EDA shows that there are violations in the assumptions as there is a multicollinearity issue due to a high correlation in

bedrooms, bathrooms, and square_feet variables.

Violations in the basic linear model assumption should be appropriately addressed. The multicollinearity violation can be resolved by quantifying multicollinearity using **Variance Inflation Factor** and dropping the input variables with VIF greater than 5.

```
[11]: # Resolve multicollinearity using vif function and remove the input variable
      ↪with vif > 5
input <- apartmentsTidy %>%
  select(-c("region", "pets_allowed"))
MLR_apartments <- lm(price ~ ., data = input)
vif_MLR_apartments <- vif(MLR_apartments)

round(vif_MLR_apartments, 3)
```

bathrooms 2.459 bedrooms 2.305 square_feet 2.939 no_of_amenities 1.013

As the VIF values are all less than 5, multicollinearity will not be an issue.

```
[12]: # Set seed to generate a reproducible random sample
      set.seed(634)

      # Split data into training and testing data
      apartmentsTidy$ID <- rownames(apartmentsTidy)
      training <- apartmentsTidy %>%
        group_by(region, pets_allowed) %>%
        sample_frac(0.75)
      testing <- anti_join(apartmentsTidy,
        training,
        by = join_by(ID)
      )

      # Remove "ID"
      training <- training %>% select(-c("ID"))
      testing <- testing %>% select(-c("ID"))
```

```
[13]: # Apply forward stepwise selection algorithm for prediction model selection and
      ↪maximum subset size can be 11
forward_sel <- regsubsets(
  x = price ~ ., nvmax = 11,
  data = training,
  method = "forward"
)

# Summarise results of algorithm
forward_sel_summary <- summary(forward_sel)

forward_sel_summary <- tibble(
  n_input_variables = 1:10,
```

```

RSS = forward_sel_summary$rss,
BIC = forward_sel_summary$bic,
Cp = forward_sel_summary$cp
)

```

```

[14]: # Forward Selection Summary
forward_sel_summary

```

A tibble: 10 × 4

n_input_variables <int>	RSS <dbl>	BIC <dbl>	Cp <dbl>
1	1067877936	-1002.368	1747.798962
2	905296331	-2097.203	466.993254
3	855670006	-2465.111	77.430301
4	848174835	-2515.092	20.291702
5	846197617	-2521.879	6.691003
6	845772858	-2516.427	5.339553
7	845603911	-2508.955	6.006525
8	845519534	-2500.815	7.340771
9	845492720	-2492.219	9.129202
10	845476345	-2483.541	11.000000

Table 5. Forward Stepwise Selection Algorithm Summary

```

[15]: # Plot Mallory's Cp for number of input variables
plot(summary(forward_sel)$cp,
main = "Cp for forward selection",
xlab = "Number of Input Variables", ylab = "Rsqr", type = "b", pch = 19,
col = "red"
)

```

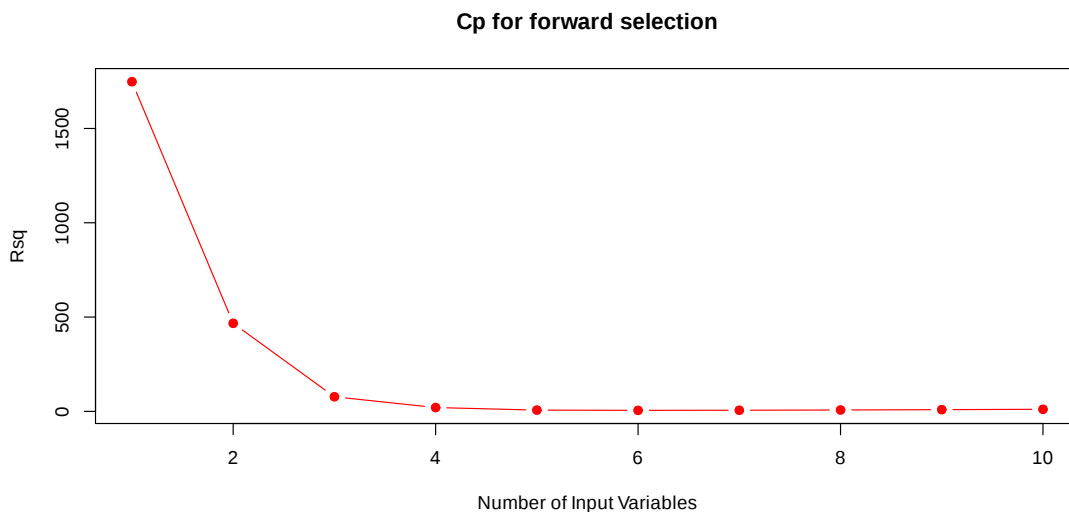


Figure 5. Line Plot Mallory's C_p vs the Number of Input variables

```
[16]: summary(forward_sel)
```

Subset selection object

Call: regsubsets.formula(x = price ~ ., nvmax = 11, data = training,
method = "forward")

10 Variables (and intercept)

	Forced in	Forced out
bathrooms	FALSE	FALSE
bedrooms	FALSE	FALSE
pets_allowedNone	FALSE	FALSE
pets_allowedOnlyCats	FALSE	FALSE
pets_allowedOnlyDogs	FALSE	FALSE
pets_allowedYes	FALSE	FALSE
square_feet	FALSE	FALSE
no_of_amenities	FALSE	FALSE
regionEastCoast	FALSE	FALSE
regionWestCoast	FALSE	FALSE

1 subsets of each size up to 10

Selection Algorithm: forward

		bathrooms	bedrooms	pets_allowedNone	pets_allowedOnlyCats
1	(1)	" "	" "	" "	" "
2	(1)	" "	" "	" "	" "
3	(1)	" "	" "	" "	" "
4	(1)	"*"	" "	" "	" "
5	(1)	"*"	" "	" "	" "
6	(1)	"*"	" "	" "	" "
7	(1)	"*"	"*"	" "	" "
8	(1)	"*"	"*"	" "	" "
9	(1)	"*"	"*"	" "	"*"
10	(1)	"*"	"*"	"*"	"*"

		pets_allowedOnlyDogs	pets_allowedYes	square_feet	no_of_amenities
1	(1)	" "	" "	"*"	" "
2	(1)	" "	" "	"*"	" "
3	(1)	" "	" "	"*"	" "
4	(1)	" "	" "	"*"	" "
5	(1)	" "	"*"	"*"	" "
6	(1)	"*"	"*"	"*"	" "
7	(1)	"*"	"*"	"*"	" "
8	(1)	"*"	"*"	"*"	"*"
9	(1)	"*"	"*"	"*"	"*"
10	(1)	"*"	"*"	"*"	"*"

		regionEastCoast	regionWestCoast
1	(1)	" "	" "
2	(1)	" "	"*"
3	(1)	"*"	"*"
4	(1)	"*"	"*"


```

5 ( 1 ) "*"          "*"
6 ( 1 ) "*"          "*"
7 ( 1 ) "*"          "*"
8 ( 1 ) "*"          "*"
9 ( 1 ) "*"          "*"
10 ( 1 ) "*"         "*"

```

Choosing the number of predictors associated with the minimum Mallory's C_p aims to strike a balance between model simplicity and adequate model performance.

```

[17]: # Choose variables for best performance by choosing the variables that
      ↪ minimizes Cp
cp_min = which.min(forward_sel_summary$Cp)
selected_var <- names(coef(forward_sel, cp_min))[-1]
selected_var

```

1. 'bathrooms' 2. 'pets_allowedOnlyDogs' 3. 'pets_allowedYes' 4. 'square_feet' 5. 'regionEastCoast' 6. 'regionWestCoast'

```

[18]: # create subset of the training data from the selected variables
encoded_data <- model.matrix(~ . - 1 + as.factor(pets_allowed) + as.
      ↪ factor(region), data = training)
df_train <- as.data.frame(encoded_data)
training_subset <- df_train %>% select(all_of(selected_var),price)

```

```

[19]: # create subset of the testing data from the selected variables
encoded_test <- model.matrix(~ . - 1 + as.factor(pets_allowed) + as.
      ↪ factor(region), data = testing)
df_test <- as.data.frame(encoded_test)
testing_subset <- df_test %>% select(all_of(selected_var),price)

```

```

[20]: # Use selected variables to make a Multiple Linear regression model
red_apartments_OLS <- lm(price ~ ., training_subset)
summary(red_apartments_OLS)

```

Call:

```
lm(formula = price ~ ., data = training_subset)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-1486.97	-250.98	-48.38	208.60	1350.02

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	703.47174	12.42554	56.615	< 2e-16 ***
bathrooms	95.25341	12.61228	7.552	4.84e-14 ***
pets_allowedOnlyDogs	82.01423	44.79389	1.831	0.0672 .

```

pets_allowedYes      47.46236    11.76405    4.035 5.53e-05 ***
square_feet          0.33555     0.01678   19.992 < 2e-16 ***
regionEastCoast      184.82219   10.23534   18.057 < 2e-16 ***
regionWestCoast      521.28899   13.47246   38.693 < 2e-16 ***
---

```

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 356 on 6675 degrees of freedom

Multiple R-squared: 0.3201, Adjusted R-squared: 0.3195

F-statistic: 523.8 on 6 and 6675 DF, p-value: < 2.2e-16

```

[21]: # Predict testing data using the reduced model
test_red_pred_OLS <- predict(red_apartments_OLS, testing_subset)
head(test_red_pred_OLS)

# Predict testing data using the full model
MLR_full_OLS <- lm(price ~ ., training)
test_full_pred_OLS <- predict(MLR_full_OLS, testing)
head(test_full_pred_OLS)

```

```

1    1019.45166020069 2    1364.30732196989 3    844.360553229618 4    847.716097861434 5
858.453840683246 6    1383.76948083443

1    1026.85734447088 2    1370.98700088722 3    852.110577592813 4    855.554717527098 5
866.575965316808 6    1389.81775123241

```

```

[22]: # Function to calculate RMSE
rmse <- function(actual, predicted) {
  sqrt(mean((actual - predicted)^2))
}

# Compare results of RMSE of the reduced and full model to see if the reduced
  ↪ model is better
results <- rbind(tibble(
  Model = "OLS Full Regression",
  RMSE = rmse(testing$price, test_full_pred_OLS)),
  tibble(
    Model = "OLS Reduced Regression using Forward Selection",
    RMSE = rmse(testing_subset$price, test_red_pred_OLS)))
results

```

	Model	RMSE
	<chr>	<dbl>
A tibble: 2 × 2	OLS Full Regression	335.1402
	OLS Reduced Regression using Forward Selection	335.0687

Table 6. RMSE for OLS Full and Reduced Regression model

Observations: The reduced OLS Regression Model is made using variables `bathrooms`, `square_feet`, and the subcategories of the `region` variable (`regionWestCoast` and `regionEastCoast`) and `pets_allowed` variable (`pets_allowedOnlyDogs` and `pets_allowedYes`) and the RMSE observed can be interpreted as the average error in predicting the apartment rental price using the listed parameters is 335.0687 USD. This is fairly reasonable as the average prices range from 200 USD to 2400 USD. The model also shows that the major apartment price influencers are appropriate variables for the predictive model selection.

0.4 Discussion

Summary, and their Implications/Impact of the Results The reduction in the number of variables from 10 to 6 variables, which creates a simpler model without losing prediction performance. Table 6 suggests that the reduced model is slightly more efficient and accurate in predicting the apartment rental price compared to the full-model. This could imply that the selected variables included in the reduced model are the most significant predictors of apartment rental prices.

Model Variable Interpretation: - `bathrooms`: Keeping other variables constant, a unit increase in the number of bathrooms increases the apartment price by 95.25 USD. - `pets_allowedOnlyDogs`: For “Only Dogs” pets policy, the apartment prices are 82.01 USD more on average than the apartments with policies other than “Only Dogs” and “No restriction” pets policy, holding other variables constant. - `square_feet`: Assuming all other variables are unchanged, a unit increase in the square footage of the apartment increases the apartment price by 0.34 USD. - `pets_allowedYes`: For the “No Restriction” pets policy, the apartment prices are 47.46 USD more on average than the apartments with policies other than “Only Dogs” and “No restriction” pets policy, holding other variables constant. - `regionWestCoast`: For Apartments in the “West Coast” Region, the apartment prices are 521.29 USD more on average than the apartments located in the “Central Midwest” Region, holding other variables constant. - `regionEastCoast`: For Apartments in the “East Coast” Region, the apartment prices are 184.82 USD more on average than the apartments located in the “Central Midwest” Region, holding other variables constant.

These insights are beneficial to both landlords and renters as they can forecast the demanding apartment trend so landlords can adjust renting prices and it also serves as the assessment to match renter’s preferences for best budgeting, and it can provide a price range of rents based on renters’ preferences.

Expectations vs Results We were expecting the full model to overfit the training data resulting in a greater variation from the actual prices in the testing data. However, the reduced model performance is just slightly better. The difference in predictive performance is very small, it suggests that the additional variables in the full model do not significantly improve the predicting performance but rather overfit the training data.

We also expect that bathrooms will be included in the reduced model though the correlation between `bathrooms` and apartment price is small, yet, there is a logical expectation that apartments with more bathrooms tend to have higher rental prices.

In Figure2, we observe that `bedrooms` can be selected while the `no_of_amenities` is not a good choice, but the forward selection in the method does not lead us to select the bedrooms. This may be because the number of bedrooms is highly correlated with the number of bathrooms and apartment

size, causing the forward selection to prioritise the other variables to avoid the multicollinearity problem.

We were also expecting multicollinearity between bedrooms, bathrooms and square_feet to be closely correlated intuitively, however, after testing the VIF, we have no issue regarding the multicollinearity between variables.

How can the Model be Improved? The proposed model can predict the price of apartments using multiple linear regression to some extent but the accuracy is still limited. Though, is is not the most optimised prediction model available, we seek to explore other choices to find one with minimed Mallows's C_p , AIC and BIC.

We also want to use K-Folds cross validation to train the model more efficiently. K-Fold can be folded multiple times to obtain more reliable estimates of model performance than splits that rely on individual training tests. and helps to mitigate the effects of randomly partitioning data into training and test sets.

Future Questions or Research from this Study Housing price is a controversial topic that affects many people's lives. In our study, we used various predictors to estimate the apartment price in the US. However, our results do not show how much each of the selected predictors contributes to the apartment price. The research can serve as a baseline for future studies that aim to understand the causal relationship between different factors and housing prices. By identifying the key factors that influence housing prices, our research also provides valuable insights for policymakers who want to regulate the apartment market more effectively.

0.5 References

- Apartment for rent classified. (2019). UCI Machine Learning Repository. <https://doi.org/10.24432/C5X623>.
- Wickramaarachchi, Nishani. (2016). Determinants of rental value for residential properties: A land owner's perspective for boarding homes. Built-Environment Sri Lanka. [link](#).
- Collinson, Rob and Winter, Ben .(2009). U.S. Rental Housing Characteristics: Supply, Vacancy, and Affordability. Available at SSRN: [link](#)